

CS 5350/6350: Machine Learning Fall 2016

Homework 4

Handed out: Oct 18, 2016

Due date: Nov 1, 2016

General Instructions

- You are welcome to talk to other members of the class about the homework. I am more concerned that you understand the underlying concepts. However, you should write down your own solution. Please keep the class collaboration policy in mind.
- Feel free discuss the homework with the instructor or the TAs.
- Your written solutions should be brief and clear. You need to show your work, not just the final answer, but you do *not* need to write it in gory detail. Your assignment should be **no more than 10 pages**. Every extra page will cost a point.
- Handwritten solutions will not be accepted.
- The homework is due by midnight of the due date. Please submit the homework on Canvas.
- Some questions are marked **For 6350 students**. Students who are registered for CS 6350 should do these questions. Of course, if you are registered for CS 5350, you are welcome to do the question too, but you will not get any credit for it.

Note

Do not just put down an answer. We want an explanation. No points will be given for just a statement of the results of a proof. You will be graded on your reasoning, not just on your final result.

Please follow good proof technique; what this means is if you make assumptions, state them. If what you do between one step and the next is not trivial or obvious, then state how and why you are doing what you are doing. A good rule of thumb is if you have to ask yourself whether what you're doing is obvious, then it's probably not obvious. Try to make the proof clean and easy to follow.

1 PAC learning

1. [20 points total] A factory assembles a product that consist of different parts. Suppose a robot was invented to recognize whether a product contains all the right parts. The rules of making products are very simple: 1) you are free to combine any of the parts

as they are 2) you may also cut any of the parts into two distinct pieces before using them. You wonder how much effort a robot would need to figure out the what parts are used in the product.

- (a) [5 points] Suppose that a naive robot has to recognize products made using only rule 1. Given N available parts and each product made out of these constitutes a distinct hypothesis. How large would the hypothesis space be? Brief explain your answer.
 - (b) [5 points] Suppose that an experienced worker follows both rules when making a product. How large is the hypothesis space now? Explain.
 - (c) [10 points] An experienced worker decides to train the naive robot to discern the makeup of a product by showing you the product samples he has assembled. There are 6 available parts. If the robot would like to learn any product at 0.01 error with probability 99%, how many examples would the robot have to see?
2. [20 points, from Tom Mitchell's book] We have learned an expression for the number of training examples sufficient to ensure that every hypothesis will have true error no worse than ϵ plus its observed training error $error_S(h)$. In particular, we used Hoeffding bounds to derive

$$m \geq \frac{1}{2\epsilon^2}(\ln(|H|) + \ln(1/\delta)).$$

Derive an alternative expression for the number of training examples sufficient to ensure that every hypothesis will have true error no worse than $(1 + \epsilon)error_S(h)$, where $0 \leq \epsilon \leq 1$. You can use general Chernoff bounds to derive such a result.

Chernoff bounds: Suppose X_1, \dots, X_m are the outcomes of m independent coin flips (Bernoulli trials), where the probability of heads on any single trail is $Pr[X_i = 1] = p$ and the probability of tails is $Pr[X_i = 0] = 1 - p$. Define $S = X_1 + X_2 + \dots + X_m$ to be the sum of these m trials. The expected value of S/m is $E[S/m] = p$. The Chernoff bounds govern the probability that S/m will differ from p by some factor $0 \leq \gamma \leq 1$.

$$\begin{aligned} Pr[S/m > (1 + \gamma)p] &\leq e^{-mp\gamma^2/3} \\ Pr[S/m < (1 - \gamma)p] &\leq e^{-mp\gamma^2/2} \end{aligned} \tag{1}$$

2 VC Dimensions

1. [10 points] Suppose you have a finite hypothesis space \mathcal{C} . Show that its VC dimension at most $\log_2 |\mathcal{C}|$ (Hint: You also prove this by contradiction.)
2. [10 points] Given some finite domain set, \mathcal{X} , and a number $k \leq |\mathcal{X}|$, figure out the VC-dimension of each of the following classes and prove your claims:
 - (a) $\mathcal{H}_{=k}^{\mathcal{X}} = \{h \in \{0, 1\}^{\mathcal{X}} : |x : h(x) = 1| = k\}$. That is , the set of all functions that assign the value 1 to exactly k elements of \mathcal{X} .

- (b) $\mathcal{H}_{\leq k}^{\mathcal{X}} = \{h \in \{0,1\}^{\mathcal{X}} : |x : h(x) = 1| \leq k \text{ or } |x : h(x) = 0| \leq k\}$. That is, the set of all functions that assign the value 1 or 0 to at most k elements of \mathcal{X} .
3. [10 points] Suppose we have an instance space consisting of real numbers and a hypothesis space \mathcal{H} consisting of *two* disjoint intervals, defined by $[a, b]$ and $[c, d]$. That is, a point $x \in \mathbb{R}$ is labeled as positive if, and only if, either $a \leq x \leq b$ or $c \leq x \leq d$. Determine the VC dimension of \mathcal{H} ?
4. [15 points] We have a learning problem where each example is a point in \mathbb{R}^2 . The concept class H is defined as follows: A function $h \in H$ is specified by two parameters a and b . An example $\mathbf{x} = \{x_1, x_2\}$ in \mathbb{R}^2 is labeled as $+$ if and only if $x_1 + x_2 \geq a$ and $x_1 - x_2 \leq b$ and is labeled $-$ otherwise.

For example, if we set $a = 0, b = 0$, the grey region in figure 1 is the region of $\mathbf{x} = \{x_1, x_2\}$ that has label $+1$.




Figure 1: An example with $a = 0, b = 0$. All points in the gray region (extending infinitely) shows the region that will be labeled as positive.

What is the VC dimension of this class?

5. [**For 6350 Students**, 15 points] Let two hypothesis classes H_1 and H_2 satisfy $H_1 \subseteq H_2$. Prove: $VC(H_1) \leq VC(H_2)$.

3 AdaBoost

[15 points] You are given the following examples in the Table 1. You need to learning a model that minimize the error on this small dataset.

Table 1: training set

$\mathbf{x} = [x_1, x_2]$	y
[1,1]	-1
[1,-1]	+1
[-1,-1]	-1
[-1,1]	-1

Assuming you are also given the following 4 weak hypothesis classifiers

$$\begin{aligned}
h_a(\mathbf{x}) &= \text{sgn}(x_1) \\
h_b(\mathbf{x}) &= \text{sgn}(x_1 - 2) \\
h_c(\mathbf{x}) &= -\text{sgn}(x_1) \\
h_d(\mathbf{x}) &= -\text{sgn}(x_2)
\end{aligned}$$

Treat them as your weak classifiers(rule of thumb) for the following question.

Step through the full AdaBoost algorithm (Lecture Boosting slide P34) for 4 rounds by choosing h_t from the above 4 weak classifiers. Remember that you need to **choose a hypothesis** from h_a, h_b, h_c, h_d whose weighted classification error is **better than chance**. However, in this question, for easier grading, we have chosen h_a as the first hypothesis and show the values of $\epsilon_1, \alpha_1, Z_1, D_1$ in Table 2.

For you answer, please follow the table template, report the hypothesis you choose and all the $\epsilon_t, \alpha_t, Z_t, D_t$, and the final hypothesis $H_{final}(x)$ for *four subsequent rounds*.

Table 2: Choose $h_a(\mathbf{x}) = \text{sgn}(x_1), \epsilon_1 = 1/4, \alpha_1 = \frac{\ln 3}{2}, Z_1 = \frac{\sqrt{2}}{3}$

$\mathbf{x} = [x_1, x_2]$	y_i	D_1	$D_1(i)y_i h_t(\mathbf{x}_i)$	D_2
[1,1]	-1	1/4	-1/4	
[1,-1]	+1	1/4	1/4	
[-1,-1]	-1	1/4	1/4	
[-1,1]	-1	1/4	1/4	